

Implementation Tutorial : Image Generation using GAN

Jin Myung Kwak

MLAI, KAIST

2020/07/27

Overview

In this tutorial, we will learn how to implement:

Image Generation using GAN

- Deep Convolutional Generative Adversarial Networks (**DCGAN**)

*Both the codes and the dataset for this tutorial will be provided by the instructor.

*The provided version may have been *slightly* modified from the original codes.

Environments

Prerequisites

- Linux or macOS
- Python ≥ 3.7
- Tensorflow = 2.2.0

Github Repositories

- `git clone https://github.com/jin8/gan.git`

Image Generation using GAN (***DCGAN***)

Download Tutorial

You can download the DCGAN code as below:

```
myung@ai2:/st2/myung/tutorial/$ git clone https://github.com/jin8/gan.git

Cloning into 'gan'...
remote: Enumerating objects: 9, done.
remote: Counting objects: 100% (9/9), done.
remote: Compressing objects: 100% (7/7), done.
remote: Total 9 (delta 1), reused 8 (delta 0), pack-reused 0
Unpacking objects: 100% (9/9), done.
```

This is what you will see after git clone.


```
- gan/
  └─ dcgan.py
  └─ dcgan_sol.py
```

Datasets

Two popular image datasets



MNIST



Fashion MNIST

How to Execute the Code?

You can run the program with this command:

```
myung@ai2:/st2/myung/tutorial/gan$ python dcgan_sol.py --dataset mnist
```

Once you execute the code you will see new folders created in the directory

```
-gan/  
├ dcgan.py  
├ images/  
│   ├── dcgan_sol_mnist/  
│   │   ├── image_at_epoch_0.png  
│   │   ├── image_at_epoch_1.png  
│   │   ├── image_at_epoch_2.png  
│   │   └── ...  
├ logs  
└ dcgan_sol_mnist.gif
```

Code Structure

We need look no further than just 1 file: *dcgan.py*

```
1. import packages
2. parser.add_argument()
3. Dataset
4. Models
   - def generator_model():
   - def discriminator_model():
5. Loss functions
   - def generator_loss():
   - def discriminator_loss():
6. Optimizer & Checkpoint
7. Training!
   for epoch in range(args.epoch):
       for i, data in enumerate(dataset):
           # train!
```


Code Structure

We need look no further than just 1 file: *dcgan.py*

```
1. import packages
2. parser.add_argument()
3. Dataset
4. Models
   - def generator_model():
   - def discriminator_model():
5. Loss functions
   - def generator_loss():
   - def discriminator_loss():
6. Optimizer & Checkpoint
7. Training!
   for epoch in range(args.epoch):
       for i, data in enumerate(dataset):
           # train!
```

1. Import Packages

This are the following packages used to train the GAN

```
import os
import time
import argparse
import tensorflow as tf
from tensorflow.keras import layers
from utils import generate_and_save_images, generate_and_save_gif

# assign specific gpu
os.environ["CUDA_VISIBLE_DEVICES"] = "#"
```

- We have installed all the packages you'll need for this project
- Please refer environment.yml if you want to check which packages are used for this tutorial
- Double check your `CUDA_VISIBLE_DEVICES` number

2. Define Argparser

```
parser = argparse.ArgumentParser()
parser.add_argument('--dataset', required=True, help='mnist | fashion')
...
parser.add_argument('--image_dir', type=str, default='./images',
                    help='directory for output images')
parser.add_argument('--log_dir', type=str, default='./logs',
                    help='directory for logging losses using tensorboard')
...
parser.add_argument('--epochs', type=int, default=100,
                    help='training epochs')
parser.add_argument('--batch_size', type=int, default=64, help='input batch size')
parser.add_argument('--latent_size', type=int, default=100,
                    help='size of the latent z vector')
```

We can assign parameters in command line:

```
myung@ai2:/st2/myung/tutorial/gan$ python dcgan.py --dataset celeba
```

3. Dataset Loading & Preprocessing


Datasets are load depending on the `args.dataset`

```
if args.dataset == 'mnist':
    (train_images, train_labels), (_, _) = tf.keras.datasets.mnist.load_data()


elif args.dataset == 'fashion':
    (train_images, train_labels), (_, _) = tf.keras.datasets.fashion_mnist.load_data()
    ...

train_images = train_images.reshape(
    train_images.shape[0],
    IMAGE_SIZE, IMAGE_SIZE,
    IMAGE_CHANNEL).astype('float32')
train_images = (train_images - 127.5) / 127.5
train_dataset = tf.data.Dataset.from_tensor_slices(train_images)\
    .shuffle(train_images.shape[0], reshuffle_each_iteration=True)\
    .batch(BATCH_SIZE)
```

4. Model : Deep Convolutional GAN (*DCGAN*)




4. Implement `def generator_model():`




- `layers.Dense(units, use_bias=False, input_shape=(input_dim,))`
- `layers.Reshape(target_shape)`
- `layers.Conv2DTranspose(filters, kernel_size, strides=(1, 1), padding='same', use_bias=False)`
- `layers.BatchNormalization()`
- `layers.LeakyReLU()`

4. Implement `def generator_model():`



- `layers.Dense(units, use_bias=False, input_shape=(input_dim,))`
- `layers.Reshape(target_shape)`
- `layers.Conv2DTranspose(filters, kernel_size, strides=(1, 1), padding='same', use_bias=False)`
- `layers.BatchNormalization()`
- `layers.LeakyReLU()`

4. Implement `def generator_model():`



BatchNorm

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_1 \dots x_m\}$;

Parameters to be learned: γ, β

Output: $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$



$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$


- `layers.Dense(units, use_bias=False, input_shape=(input_dim,))`
- `layers.Reshape(target_shape)`
- `layers.Conv2DTranspose(filters, kernel_size, strides=(1, 1), padding='same', use_bias=False)`
- `layers.BatchNormalization()`
- `layers.LeakyReLU()`

4. Implement `def generator_model():`





- `layers.Dense(units, use_bias=False, input_shape=(input_dim,))`
- `layers.Reshape(target_shape)`
- `layers.Conv2DTranspose(filters, kernel_size, strides=(1, 1), padding='same', use_bias=False)`
- `layers.BatchNormalization()`
- `layers.LeakyReLU()`

4. Implement `def generator_model():`




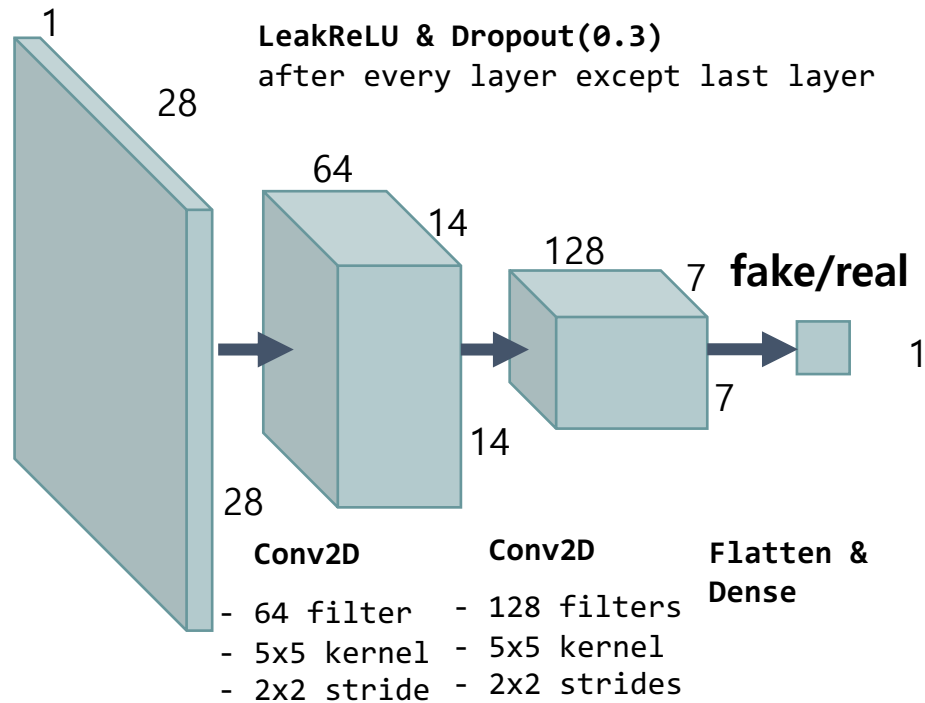
- `layers.Dense(units, use_bias=False, input_shape=(input_dim,))`
- `layers.Reshape(target_shape)`
- `layers.Conv2DTranspose(filters, kernel_size, strides=(1, 1), padding='same', use_bias=False)`
- `layers.BatchNormalization()`
- `layers.LeakyReLU()`

4. Implement `def discriminator_model():`





- `layers.Conv2D(filters, kernel_size, strides=(1, 1), padding='same', use_bias=False)`
- `layers.LeakyReLU()`
- `layers.Dropout(rate)`
- `layers.Flatten()`
- `layers.Dense(units, use_bias=False, input_shape=(input_dim,))`

4. Implement def discriminator_model():



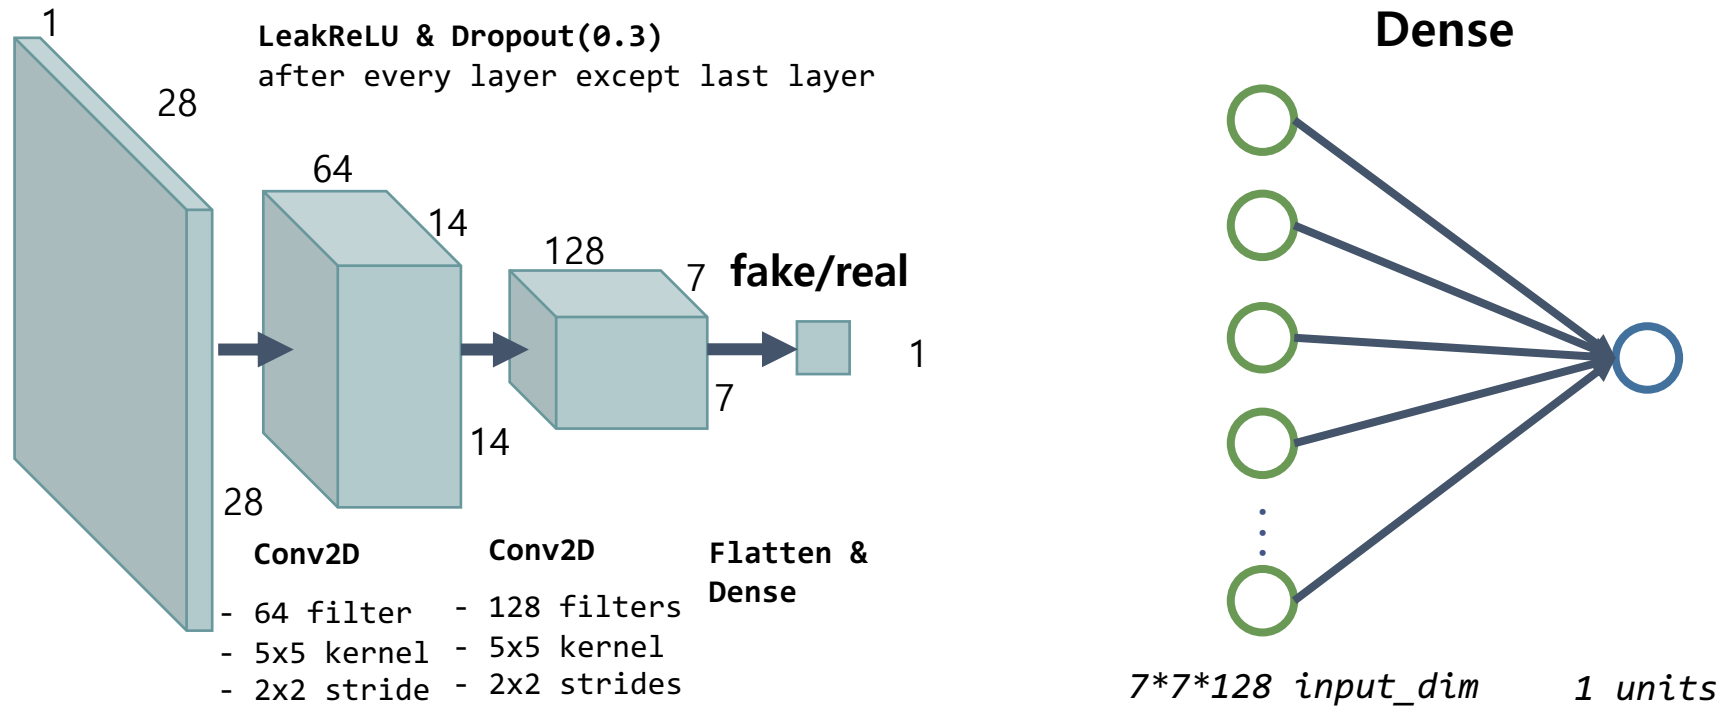
- `layers.Conv2D(filters, kernel_size, strides=(1, 1), padding='same', use_bias=False)`
- `layers.LeakyReLU()`
- `layers.Dropout(rate)`
- `layers.Flatten()`
- `layers.Dense(units, use_bias=False, input_shape=(input_dim,))`

4. Implement `def discriminator_model():`



- `layers.Conv2D(filters, kernel_size, strides=(1, 1), padding='same', use_bias=False)`
- `layers.LeakyReLU()`
- `layers.Dropout(rate)`
- `layers.Flatten()`
- `layers.Dense(units, use_bias=False, input_shape=(input_dim,))`

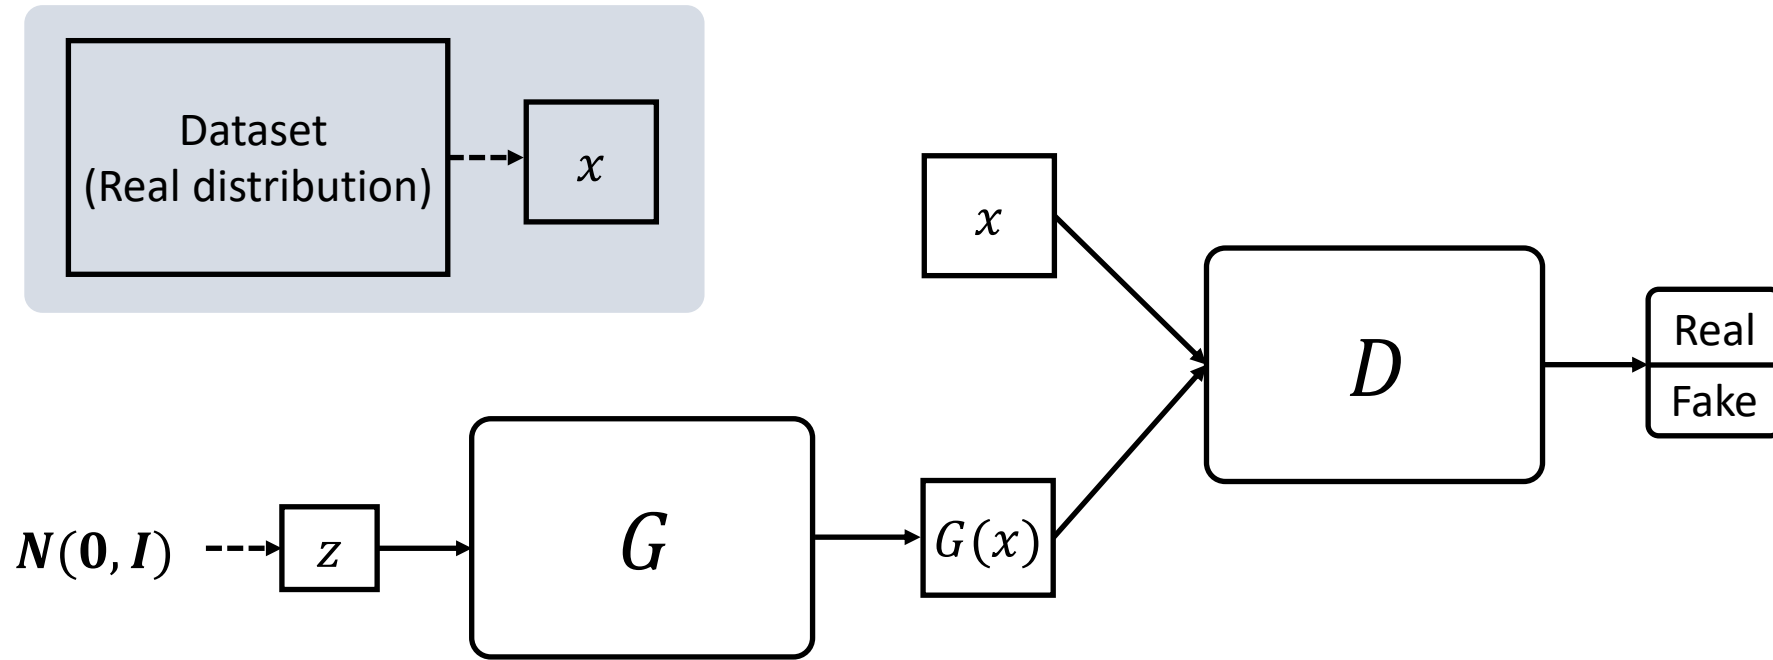
4. Implement `def discriminator_model():`



- `layers.Conv2D(filters, kernel_size, strides=(1, 1), padding='same', use_bias=False)`
- `layers.LeakyReLU()`
- `layers.Dropout(rate)`
- `layers.Flatten()`
- `layers.Dense(units, use_bias=False, input_shape=(input_dim,))`

5. Loss Function

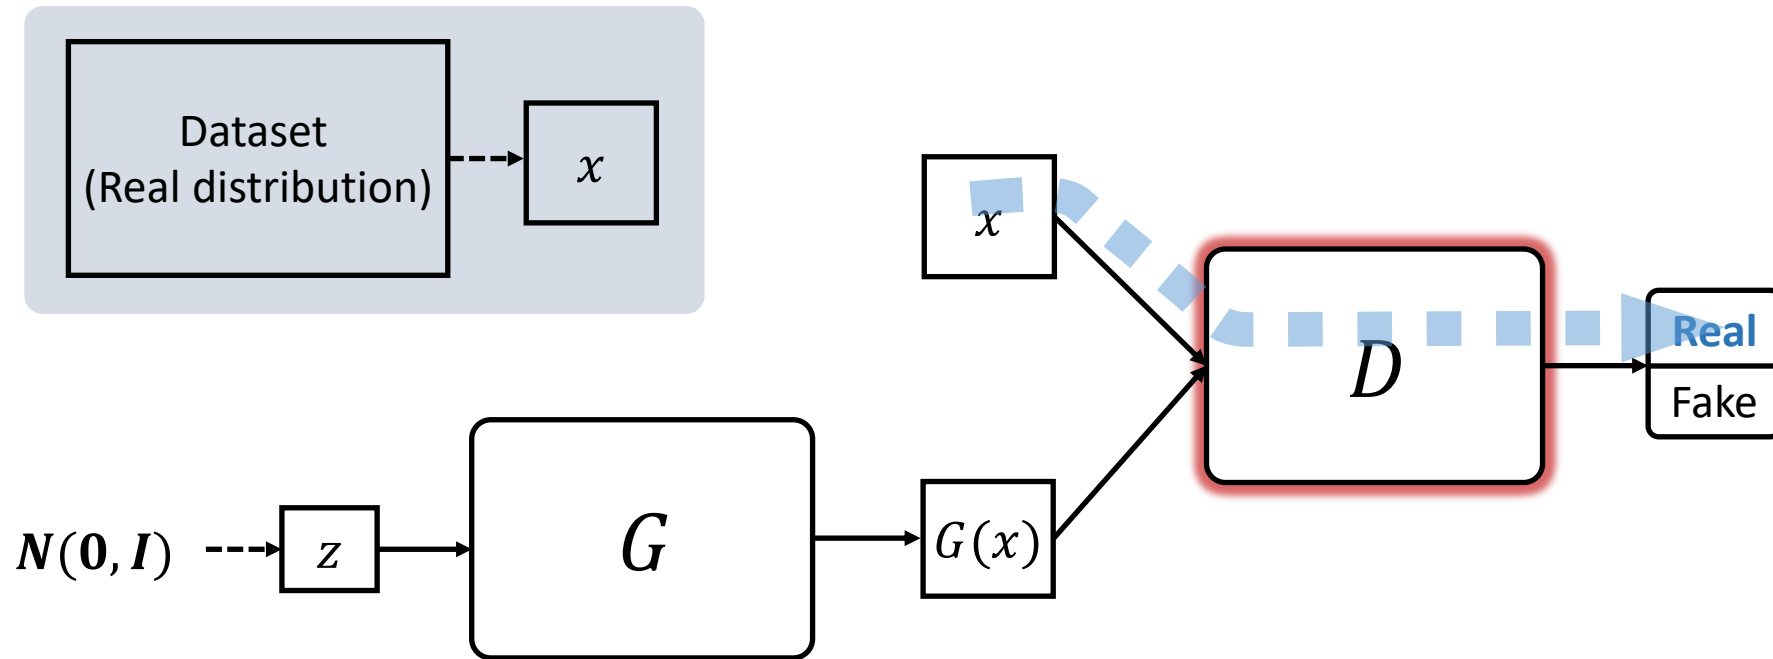
Recap: Overview



$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))].$$

5. Loss Function


Recap: Training of *Discriminator* ($Real \rightarrow Real$)



$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))].$$

5. Loss Function

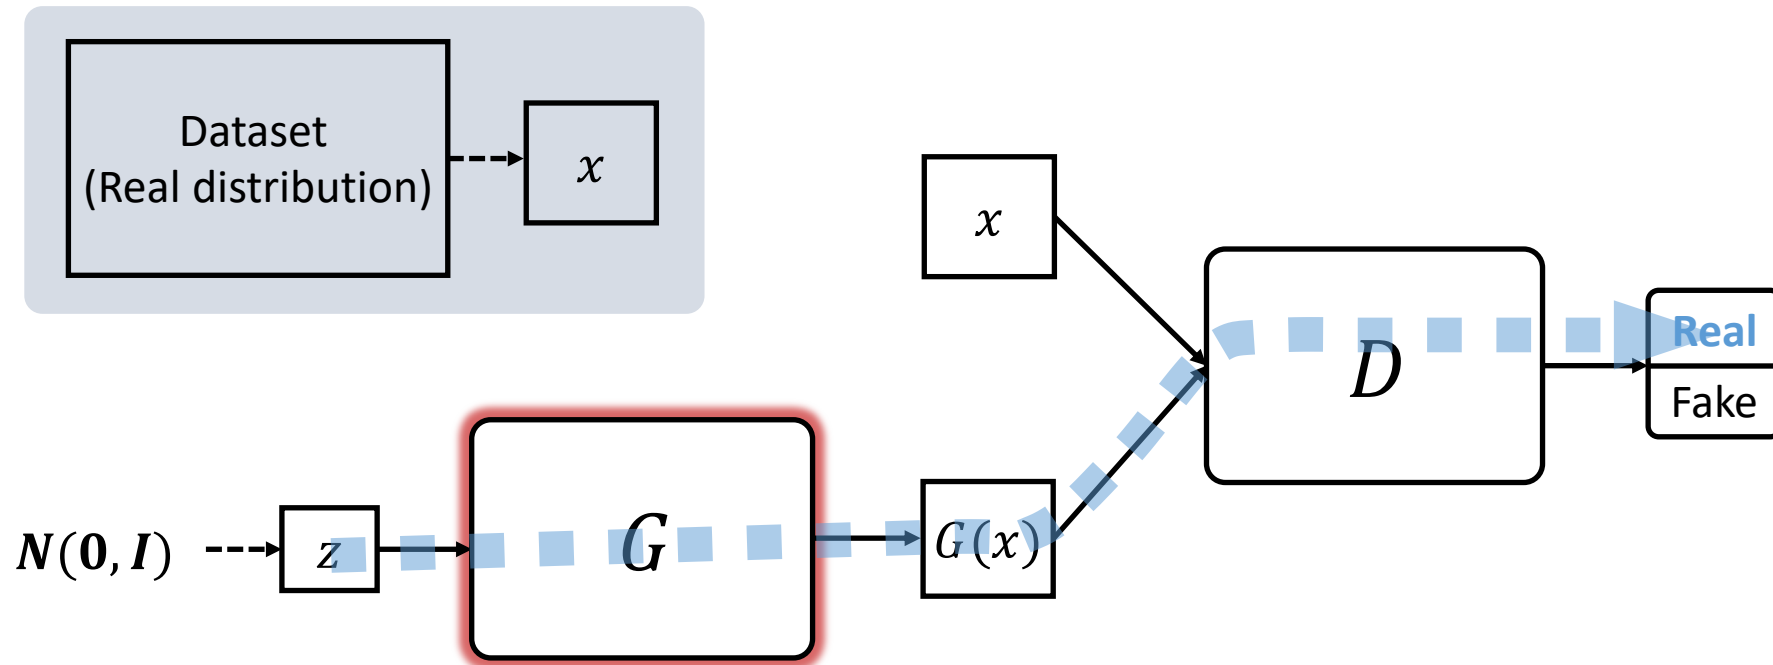
Recap: Training of **Discriminator** (*Fake* \rightarrow *Fake*)



$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))].$$

5. Loss Function

Recap: Training of **Generator** (*Fake* \rightarrow *Real*)



$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))].$$

5. Loss Function

Loss function and Optimization

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_z(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

Discriminator


prob. of D predicting that
Real data is genuine.

prob. of D predicting that
fake data is **NOT** genuine.

$$\text{Maximize : } \log(D_1(x)) + \log(1 - D_2(G(z)))$$

$\log(x)$

$\log(\text{sgm}(x))$





Generator

prob. of D predicting that
fake data is **NOT** genuine.

$$\text{Minimize : } \log(1 - D_2(G(z)))$$

$$\text{Maximize : } \log(D_2(G(z)))$$

prob. of D predicting that
fake data is genuine.



⊙: D rejecting $G(x)$
with high confidence.

5. Loss Function

Loss function and Optimization

We will use the Binary Cross Entropy loss function.

```
cross_entropy = tf.keras.losses.BinaryCrossentropy(from_logits=True)
loss = cross_entropy(y_true, y_pred)
```

We can specify which part of the equation to use with the label y .

$$l_n = - \left[\underbrace{y_n \cdot \log x_n}_{\substack{\text{We use this part,} \\ \text{when label is \textit{real}}}} + \underbrace{(1 - y_n) \cdot \log(1 - x_n)}_{\substack{\text{We use this part,} \\ \text{when label is \textit{fake}}} \right]$$

Real label: $y = 1$
Fake label: $y = 0$

6. Optimizer & Checkpoint

Since two networks are trained separately, two optimizers are defined.

```
generator_optimizer = tf.keras.optimizers.Adam(LR)
discriminator_optimizer = tf.keras.optimizers.Adam(LR)
```

We can save and restore the model using `tf.train.Checkpoint`

```
checkpoint_prefix = os.path.join(CKPT_DIR, "ckpt")
checkpoint = tf.train.Checkpoint(
    generator_optimizer=generator_optimizer,
    discriminator_optimizer=discriminator_optimizer,
    generator=generator, discriminator=discriminator)

checkpoint.save(file_prefix = checkpoint_prefix) # save model
...
checkpoint.restore(tf.train.latest_checkpoint(CKPT_DIR)) # restore model
```

7. Implement `def train_step(images):`

```
@tf.function
def train_step(images):
    # 1. sample noise z from normal distribution
    with tf.GradientTape() as gen_tape, tf.GradientTape() as disc_tape:
        # 2. generator generates fake image
        # 3. discriminator discriminates real image
        # 4. discriminator discriminates fake image
        # 5. compute discriminator loss
        # 6. compute generator loss

    gradients_of_generator = gen_tape.gradient(
        gen_loss, generator.trainable_variables)
    gradients_of_discriminator = disc_tape.gradient(
        disc_loss, discriminator.trainable_variables)

    generator_optimizer.apply_gradients(
        zip(gradients_of_generator, generator.trainable_variables))
    discriminator_optimizer.apply_gradients(
        zip(gradients_of_discriminator, discriminator.trainable_variables))
```

7. Training

```
for epoch in range(EPOCHS):
    start = time.time()

    for image_batch in train_dataset:
        train_step(image_batch)



    with train_summary_writer.as_default():
        tf.summary.scalar('gen_loss', train_gen_loss.result(), step=epoch)
        tf.summary.scalar('disc_loss', train_disc_loss.result(), step=epoch)
    generate_and_save_images(generator, epoch + 1, seed_z, IMAGE_DIR)

    if (epoch + 1) % 15 == 0:
        checkpoint.save(file_prefix = checkpoint_prefix)
        print ('Time for epoch {} is {} sec'.format(epoch + 1, time.time()-start))

generate_and_save_images(generator, EPOCHS, seed_z, IMAGE_DIR)
generate_and_save_gif(IMAGE_DIR, GIF_FILE)
```

Result

After running dcgan.py, you will have gif files.



Any questions?